


Figure 6 — Cross-provider model evaluation across three Lead-Optimizer agent stages


Stage 2: 14 vision models × 8 reps on a fixed LID. Stage 5: 11 text models × 4 project contexts × 3 reps. Stage 6: 11 text models × 4 leads × 3 reps. Bars dimmed when JSON validity < 50 % (model unreliable on this task).

openai gemini opencode mistral groq

(a) Stage 2 — Vision Agent · self-consistency Jaccard



(b) Stage 5 — Project context · endpoint-priority rubric



(c) Stage 6 — Optimization agent · SAR-strategy rubric

